

LECTURE 3

Pandas, Part 1

Introduction to Pandas syntax and operators

Sean Kang
Lecturer

Goals For This Lecture


- Introduce Pandas, with emphasis on:
 - A mental model of DataFrames - linking to statistics.
 - Key Data Structures (data frames, series, indices).
 - How to index into these structures.
 - How to read files to create these structures.
 - Other basic operations on these structures.
- Will go through quite a lot of the language without full explanations.
 - We expect you to fill in the gaps on homeworks, projects, and through your own experimentation.
- Solve some very basic data science problems using Jupyter/pandas.

Data Frames: a high-level, statistical perspective

The world, a statistician's view



A (statistical) population
from which we draw
samples.
Each sample has certain
features.



	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

A generic DataFrame
(from <https://arxiv.org/abs/2001.00888>)

Working with DataFrames

- Most commonly used data frame operators
 - `read_csv`
 - `head`
 - `loc`
 - `Groupby`
- Issues with Dataframes
 - Performance issues with relatively large datasets
- There is a research paper where they discuss building a new system MODIN which proposes to alleviate those performance issues
 - For more reading <https://arxiv.org/abs/2001.00888>

Pandas Data Structures: Data Frames, Series, and Indices

Pandas Data Structures

There are three fundamental data structures in pandas:

- Data Frame: 2D data tabular data.
- Series: 1D data. I usually think of it as columnar data.
- Index: A sequence or reference into row labels.

Series

Data Frame

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

```
0    Obama
1    McCain
2    Obama
3    Romney
4    Clinton
5    Trump
Name: Candidate, dtype: object
```


Index

The Relationship Between Data Frames, Series, and Indices

We can think of a Data Frame as a collection of Series that all share the same Index.

- Candidate, Party, %, Year, and Result Series all share an index from 0 to 5.

Candidate Series Party Series % Series Year Series Result Series



	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

Indices Are Not Necessarily Row Numbers

Indices (a.k.a. row labels) can also:

- Be non-numeric.
- Have a name, e.g. “State”.

	Motto	Translation	Language	Date Adopted
State				
Alabama	Audemus jura nostra defendere	We dare defend our rights!	Latin	1923
Alaska	North to the future	—	English	1967
Arizona	Ditat Deus	God enriches	Latin	1863
Arkansas	Regnat populus	The people rule	Latin	1907
California	Eureka (Εὕρηκα)	I have found it	Greek	1849

Indices

The row labels that constitute an index do not have to be unique.

- Left: The index values are all unique and numeric, acting as a row number.
- Right: The index values are named and non-unique.

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

Year	Candidate	Party	%	Result
2008	Obama	Democratic	52.9	win
2008	McCain	Republican	45.7	loss
2012	Obama	Democratic	51.1	win
2012	Romney	Republican	47.2	loss
2016	Clinton	Democratic	48.2	loss
2016	Trump	Republican	46.1	win

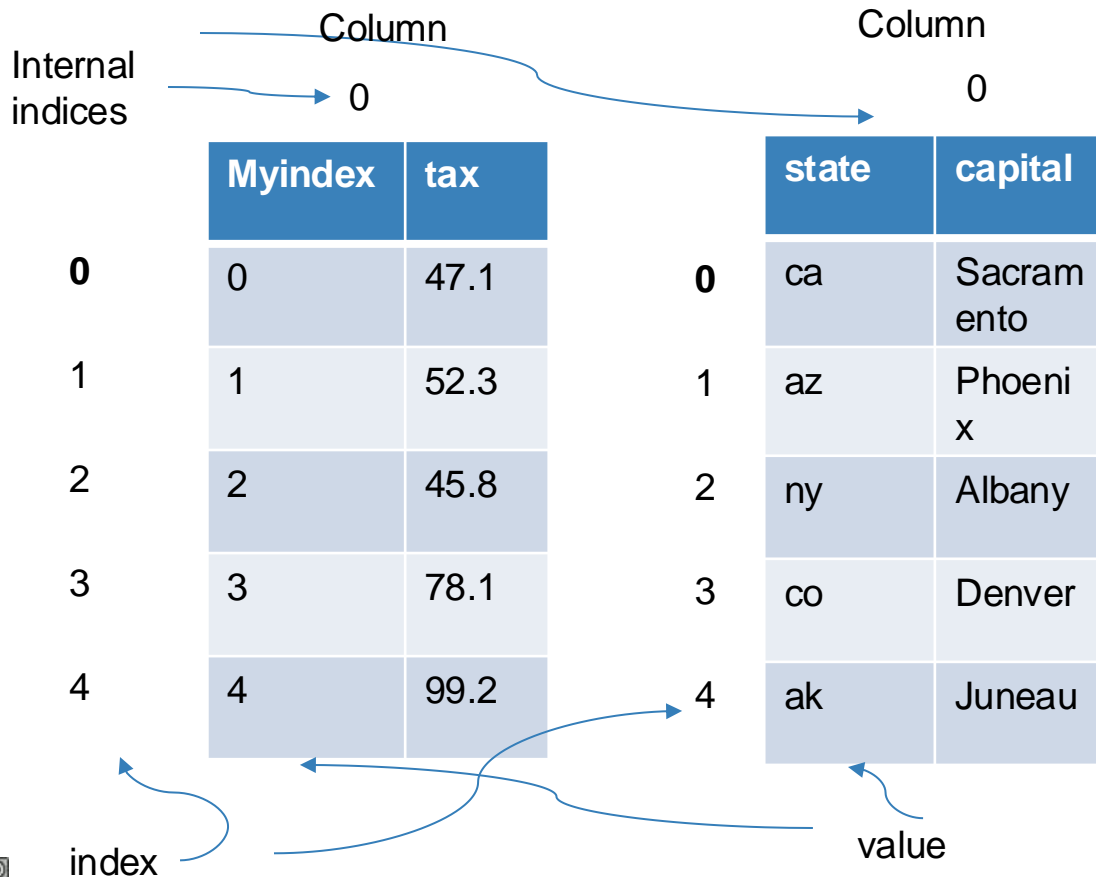
Column Names Are Usually Unique!

Column names in Pandas are almost always unique!

- Example: Really shouldn't have two columns named "Candidate".

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

Summary: structure of a Series



Access:

iLoc:
Integer/ indices

Loc:
Labels
(even if it is
integer)

Summary: structure of a DataFrame

Internal indices

0 1

Column labels

	state	Capital	Taxes
0	ca	Sacram ento	14.7
1	az	Phoenix	6.7
2	ny	Albany	7.3
3	co	Denver	12.3
4	ak	Juneau	5.6

Row labels

The diagram illustrates the structure of a DataFrame. It shows a table with 5 rows and 3 columns. The columns are labeled 'state', 'Capital', and 'Taxes'. The rows are indexed 0 to 4. Arrows point from the text 'Internal indices' to the column indices 0 and 1, and from 'Column labels' to the column headers. Another arrow points from 'Row labels' to the row index 4.

Hands On Exercise

Let's experiment with reading csv files and playing around with indices.

Show Iris dataset in Jupyter

<https://www.kaggle.com/code/jchen2186/machine-learning-with-iris-dataset>

<https://www.pycodemates.com/2022/05/iris-dataset-classification-with-python.html>

Indexing by Column Names Using [] Operator

Given a dataframe, it is common to extract a Series or a collection of Series. This process is also known as “Column Selection” or sometimes “indexing by column”.

```
elections[["Candidate"]].head(6)
```

- Column name argument to [] yields Series.
- List argument (**even of one name**) to [] yields a Data Frame.

```
elections["Candidate"].head(6)
```

```
Year
1980    Reagan
1980    Carter
1980    Anderson
1984    Reagan
1984    Mondale
1988     Bush
Name: Candidate, dtype: object
```

Candidate	
Year	
1980	Reagan
1980	Carter
1980	Anderson
1984	Reagan
1984	Mondale
1988	Bush

Indexing by Row Slices Using [] Operator


We can also index by row numbers using the [] operator.

- Numeric slice argument to [] yields rows.
- Example: [0:3] yields rows 0 to 2.

```
elections[0:3]
```

	Candidate	Party	%	Result
Year				
1980	Reagan	Republican	50.7	win
1980	Carter	Democratic	41.0	loss
1980	Anderson	Independent	6.6	loss


[] Summary



```
elections["Candidate"].head(6)
```


Year	
1980	Reagan
1980	Carter
1980	Anderson
1984	Reagan
1984	Mondale
1988	Bush

Name: Candidate, dtype: object



```
elections[["Candidate"]].head(6)
```

	Candidate
Year	
1980	Reagan
1980	Carter
1980	Anderson
1984	Reagan
1984	Mondale
1988	Bush



```
elections[0:3]
```

	Candidate	Party	%	Result
Year				
1980	Reagan	Republican	50.7	win
1980	Carter	Democratic	41.0	loss
1980	Anderson	Independent	6.6	loss

Note: Row Selection Requires Slicing!!

`elections[0]` will not work unless the elections data frame has a column whose name is the numeric zero.

- Note: It is actually possible for columns to have names that are non-String types, e.g. numeric, datetime etc.

Question

```
weird = pd.DataFrame({1: ["topdog", "botdog"], "1": ["topcat", "botcat"]})  
weird
```

	1	1
0	topdog	topcat
1	botdog	botcat

Try to predict the output of the following:

- `weird[1]`
- `weird["1"]`
- `weird[1:]`

Name → `[]` → Series
Single Column Selection

List → `[]` → DataFrame
Multiple Column Selection

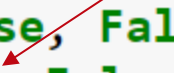
Numeric Slice → `[]` → DataFrame
(Multiple) Row Selection

Boolean Array Selection and Querying

Boolean Array Input

Yet another input type supported by [] is the boolean array.

```
elections[[False, False, False, False, False,
            False, False, True, False, False,
            True, False, False, False, True,
            False, False, False, False, False,
            False, False, True]]
```



Entry number 7

	Candidate	Party	%	Year	Result
7	Clinton	Democratic	43.0	1992	win
10	Clinton	Democratic	49.2	1996	win
14	Bush	Republican	47.9	2000	win
22	Trump	Republican	46.1	2016	win

We can generate a Boolean mask

```
[36]: iswin = elections_year_index['Result'] == 'win'  
      iswin#.head(5)
```

```
[36]: Year  
      1980      True  
      1980      False  
      1980      False  
      1984      True  
      1984      False  
      1988      True  
      1988      False  
      1992      True  
      1992      False  
      1992      False  
      1996      True  
      1996      False  
      1996      False  
      2000      False  
      2000      True  
      2004      False  
      2004      True  
      2008      True  
      2008      False  
      2012      True  
      2012      False  
      2016      False  
      2016      True  
      Name: Result, dtype: bool
```

Shows me the election year for which the result column is equal to "win"

Use the generated mask to select data

```
elections_year_index[iswin]
```

	Candidate	Party	%	Result
Year				
1980	Reagan	Republican	50.7	win
1984	Reagan	Republican	58.8	win
1988	Bush	Republican	53.4	win
1992	Clinton	Democratic	43.0	win
1996	Clinton	Democratic	49.2	win
2000	Bush	Republican	47.9	win
2004	Bush	Republican	50.7	win
2008	Obama	Democratic	52.9	win
2012	Obama	Democratic	51.1	win
2016	Trump	Republican	46.1	win

OR

```
elections_year_index[elections_year_index['Result'] == 'win']
```

Boolean Array Input

Yet another input type supported by [] is the boolean array. Useful because boolean arrays can be generated by using logical operators on Series.

Length 23 Series where every entry is "Republican", "Democrat" or "Independent."

Length 23 Series where every entry is either "True" or "False", where "True" occurs for every independent candidate.

```
elections[elections['Party'] == 'Independent']
```

	Candidate	Party	%	Year	Result
2	Anderson	Independent	6.6	1980	loss
9	Perot	Independent	18.9	1992	loss
12	Perot	Independent	8.4	1996	loss

Boolean Array Input

Boolean Series can be combined using the & operator, allowing filtering of results by multiple criteria.

```
elections[(elections['Result'] == 'win')  
          & (elections['%'] < 50)]
```

	Candidate	Party	%	Year	Result
7	Clinton	Democratic	43.0	1992	win
10	Clinton	Democratic	49.2	1996	win
14	Bush	Republican	47.9	2000	win
22	Trump	Republican	46.1	2016	win

isin

The `isin` function makes it more convenient to find rows that match one of many possible values.

Example: Suppose we want to find “Republican” or “Democratic” candidates. Could use the `|` operator (`|` means or), or we can use `isin`.

- Ugly:

```
df[(df["Party"] == "Democratic") | (df["Party"] == "Republican")]
```
- Better:

```
df[df["Party"].isin(["Republican", "Democratic"])]
```

The Query Command

The query command provides an alternate way to combine multiple conditions.

```
elections.query("Result == 'win' and Year < 2000")
```

	Candidate	Party	%	Year	Result
0	Reagan	Republican	50.7	1980	win
3	Reagan	Republican	58.8	1984	win
5	Bush	Republican	53.4	1988	win
7	Clinton	Democratic	43.0	1992	win
10	Clinton	Democratic	49.2	1996	win

Indexing with .loc and .iloc

Sampling with .sample

Loc and iloc

Loc and iloc are alternate ways to index into a DataFrame.

- They take a lot of getting used to! Documentation and ideas behind them are quite complex.
- I'll go over common usages (see docs for weirder ones).

Documentation:

- loc: <https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.loc.html>
- iloc: <https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.iloc.html>
- More general docs on indexing and selecting: [Link](#)

Loc

Loc does two things:

- Access values by labels.
- Access values using a boolean array (a la Boolean Array Selection).

Loc with Lists

The most basic use of loc is to provide a list of row and column labels, which returns a DataFrame.

```
elections.loc[[0, 1, 2, 3, 4], ['Candidate', 'Party', 'Year']]
```

	Candidate	Party	Year
0	Reagan	Republican	1980
1	Carter	Democratic	1980
2	Anderson	Independent	1980
3	Reagan	Republican	1984
4	Mondale	Democratic	1984

Loc with Lists

The most basic use of loc is to provide a list of row and column labels, which returns a DataFrame.

```
elections_year_index.loc[[1980, 1984], ['Candidate', 'Party']]
```

	Candidate	Party
Year		
1980	Reagan	Republican
1980	Carter	Democratic
1980	Anderson	Independent
1984	Reagan	Republican
1984	Mondale	Democratic

Loc with Slices

Loc is also commonly used with slices.

- Slicing works with all label types, not just numeric labels.
- Slices with loc are **inclusive**, not **exclusive**.

```
elections.loc[0:4, 'Candidate':'Year']
```

	Candidate	Party	Year
0	Reagan	Republican	1980
1	Carter	Democratic	1980
2	Anderson	Independent	1980
3	Reagan	Republican	1984
4	Mondale	Democratic	1984

Loc with Slices

Loc is also commonly used with slices.

- Slicing works with all label types, not just numeric labels.
- Slices with loc are **inclusive**, not **exclusive**.

```
elections_year_index.loc[1980:1984, 'Candidate':'Party']
```

	Candidate	Party
Year		
1980	Reagan	Republican
1980	Carter	Democratic
1980	Anderson	Independent
1984	Reagan	Republican
1984	Mondale	Democratic

Loc with Single Values for Column Label

If we provide only a single label as column argument, we get a Series.

```
elections.loc[0:4, 'Candidate']
```

```
0      Reagan
```

```
1      Carter
```

```
2    Anderson
```

```
3      Reagan
```

```
4    Mondale
```

```
Name: Candidate, dtype: object
```

Loc with Single Values for Column Label

As before with the `[]` operator, if we provide a list of only one label as an argument, we get back a dataframe.

```
elections.loc[0:4, 'Candidate']
```

0	Reagan
1	Carter
2	Anderson
3	Reagan
4	Mondale

Name: Candidate, dtype: object

```
elections.loc[0:4, ['Candidate']]
```

	Candidate
0	Reagan
1	Carter
2	Anderson
3	Reagan
4	Mondale

Loc with Single Values for Row Label

If we provide only a single row label, we get a Series.

- Such a series represents a ROW not a column!
- The index of this Series is the names of the columns from the data frame.
- Putting the single row label in a list yields a dataframe version.

```
elections.loc[0, 'Candidate':'Year']
```

Candidate	Reagan
Party	Republican
%	50.7
Year	1980

Name: 0, dtype: object

```
elections.loc[[0], 'Candidate':'Year']
```

	Candidate	Party	%	Year
0	Reagan	Republican	50.7	1980

Loc Supports Boolean Arrays

Loc supports Boolean Arrays exactly as you'd expect.

```
elections.loc[(elections['Result'] == 'win') & (elections['%'] < 50), 'Candidate':'%']
```

	Candidate	Party	%
7	Clinton	Democratic	43.0
10	Clinton	Democratic	49.2
14	Bush	Republican	47.9
22	Trump	Republican	46.1

iloc: Integer-Based Indexing for Selection by Position

In contrast to loc, iloc doesn't think about labels at all. Instead, it returns the items that appear in the numerical positions specified.

```
elections.iloc[0:3, 0:3]
```

	Candidate	Party	%
0	Reagan	Republican	50.7
1	Carter	Democratic	41.0
2	Anderson	Independent	6.6

```
mottos.iloc[0:3, 0:3]
```

	Motto	Translation	Language
State			
Alabama	Audemus jura nostra defendere	We dare defend our rights!	Latin
Alaska	North to the future	—	English
Arizona	Ditat Deus	God enriches	Latin

Advantages of loc:

- Harder to make mistakes.
- Easier to read code.
- Not vulnerable to changes to the ordering of rows/cols in raw data files.

Nonetheless, iloc can be more convenient. *Use iloc judiciously.*

Annoying Question Challenge

Which of the following pandas statements returns a DataFrame of the first 3 Candidate names only for candidates that won with more than 50% of the vote.

`elections.iloc[[0, 3, 5], [0, 3]]`

`elections.loc[[0, 3, 5], ["Candidate": "Year"]`

`elections.loc[elections["%"] > 50, ["Candidate", "Year"]].head(3)`

`elections.loc[elections["%"] > 50, ["Candidate", "Year"]].iloc[0:2, :]`

	Candidate	Party	%	Year	Result
0	Reagan	Republican	50.7	1980	win
1	Carter	Democratic	41.0	1980	loss
2	Anderson	Independent	6.6	1980	loss
3	Reagan	Republican	58.8	1984	win
4	Mondale	Democratic	37.6	1984	loss
5	Bush	Republican	53.4	1988	win
6	Dukakis	Democratic	45.6	1988	loss



	Candidate	Year
0	Reagan	1980
3	Reagan	1984
5	Bush	1988

Annoying Question Challenge

Which of the following pandas statements returns a DataFrame of the first 3 Candidate names only for candidates that won with more than 50% of the vote.

`elections.iloc[[0, 3, 5], [0, 3]]`

`elections.loc[[0, 3, 5], ["Candidate": "Year"]]`

`elections.loc[elections["%"] > 50, ["Candidate", "Year"]].head(3)`

`elections.loc[elections["%"] > 50, ["Candidate", "Year"]].iloc[0:2, :]`

	Candidate	Party	%	Year	Result
0	Reagan	Republican	50.7	1980	win
1	Carter	Democratic	41.0	1980	loss
2	Anderson	Independent	6.6	1980	loss
3	Reagan	Republican	58.8	1984	win
4	Mondale	Democratic	37.6	1984	loss
5	Bush	Republican	53.4	1988	win
6	Dukakis	Democratic	45.6	1988	loss



	Candidate	Year
0	Reagan	1980
3	Reagan	1984
5	Bush	1988

Sample

If you want a DataFrame consisting of a random selection of rows, you can use the sample method.

- By default, *it is without replacement*. Use `replace=true` for replacement.
- Naturally, can be chained with our selection operators `[]`, `loc`, `iloc`.

```
elections.sample(10)
```

	Candidate	Party	%	Year	Result
15	Kerry	Democratic	48.3	2004	loss
16	Bush	Republican	50.7	2004	win
22	Trump	Republican	46.1	2016	win
9	Perot	Independent	18.9	1992	loss
21	Clinton	Democratic	48.2	2016	loss
11	Dole	Republican	40.7	1996	loss
20	Romney	Republican	47.2	2012	loss
14	Bush	Republican	47.9	2000	win
8	Bush	Republican	37.4	1992	loss
1	Carter	Democratic	41.0	1980	loss

```
elections.query("Year < 1992").sample(4, replace=True)
```

	Candidate	Party	%	Year	Result
1	Carter	Democratic	41.0	1980	loss
4	Mondale	Democratic	37.6	1984	loss
6	Dukakis	Democratic	45.6	1988	loss
1	Carter	Democratic	41.0	1980	loss

Handy Properties and Utility Functions for Series and DataFrames

Numpy Operations

Pandas Series and DataFrames support a large number of operations, including mathematical operations so long as the data is numerical.

```
winners = elections.query("Result == 'win'")["%"]  
winners
```

0 50.7

3 58.8

5 53.4

7 43.0

10 49.2

14 47.9

16 50.7

17 52.9

19 51.1

22 46.1

Name: %, dtype: float64

```
np.mean(winners)
```

50.38

```
max(winners)
```

58.8

head, size, shape, and describe

`head`: Displays only the top few rows.

`size`: Gives the total number of data points.

`shape`: Gives the size of the data in rows and columns.

`describe`: Provides a summary of the data.

index and columns

index: Returns the index (a.k.a. row labels).

columns: Returns the labels for the columns.

The sort_values Method

One incredibly useful method for DataFrames is `sort_values`, which creates a copy of a DataFrame sorted by a specific column.

```
elections.sort_values('%', ascending=False)
```

	Candidate	Party	%	Year	Result
3	Reagan	Republican	58.8	1984	win
5	Bush	Republican	53.4	1988	win
17	Obama	Democratic	52.9	2008	win
19	Obama	Democratic	51.1	2012	win
0	Reagan	Republican	50.7	1980	win

The sort_values Method

We can also use `sort_values` on a Series, which returns a copy with the values in order.

```
mottos['Language'].sort_values().head(5)
```

State

Washington	Chinook Jargon
------------	----------------

Wyoming	English
---------	---------

New Jersey	English
------------	---------

New Hampshire	English
---------------	---------

Nevada	English
--------	---------

Name: Language, dtype: object

The value_counts Method

Series also has the function `value_counts`, which creates a new Series showing the counts of every value.

```
elections['Party'].value_counts()
```

```
Democratic      10
```

```
Republican      10
```

```
Independent       3
```

```
Name: Party, dtype: int64
```

The unique Method

Another handy method for Series is `unique`, which returns all unique values as an array.

```
mottos['Language'].unique()
```

```
array(['Latin', 'English', 'Greek', 'Hawaiian', 'Italian', 'French',  
      'Spanish', 'Chinook Jargon'], dtype=object)
```